

# **Technical Implementation Document for Complete Data Migration from RDBMS to Property Graph Model (Corrected + Enhanced)**

## **I. Project Context and Architectural Foundation**

Modern data systems face challenges due to the dramatic growth in volume, variety, and velocity of data, collectively known as Big Data.<sup>1</sup> Traditional Relational Database Management Systems (RDBMS) struggle with:

- Limited scalability
- High join costs
- Poor handling of highly interconnected data
- Fixed schema structure

To address these limitations, organizations increasingly migrate data to NoSQL systems, particularly Graph Databases, which excel at modeling complex, relationship-heavy data.<sup>1</sup>

### **A. Migration Imperative**

RDBMS systems (MySQL, Oracle, PostgreSQL) follow the relational model, store data in tables, enforce ACID guarantees, and use SQL language for DDL/DML operations.<sup>1</sup>

However, Big Data applications require:

- High concurrency
- Flexible schema
- Distributed architecture
- Fast relationship traversal

NoSQL databases were developed to address these limitations, providing schema-less, distributed, scalable data models.<sup>1</sup>

## B. Rationale for Choosing Graph Databases

Among NoSQL models—Key-Value, Document, Wide Column, Graph—the graph database is ideal when data has many interconnected entities (social networks, enterprise systems, knowledge graphs).<sup>1</sup>

Graph Databases (Neo4j):

- Use nodes, edges, and properties<sup>1</sup>
- Support fast relationship traversal<sup>1</sup>
- Remove costly “joins”<sup>1</sup>
- Provide intuitive graph modeling<sup>1</sup>
- Support Cypher, a declarative graph query language<sup>1</sup>

## C. Need for the SCT Architecture (Source → Conceptual → Target)

The research paper stresses a critical weakness:

Most existing tools use Source-to-Target (ST) translation directly, which leads to a flat and weak target schema, losing semantics like:

- Inheritance
- Aggregation
- Cardinality
- Relationship meaning/naming
- Complex constraints

Only a few studies use the Source → Conceptual → Target (SCT) path, which produces richer graph models.<sup>1</sup>

Why SCT is mandatory for this project:

- ✓ Captures implicit metadata 1
- ✓ Enhances semantics through conceptual modeling 1
- ✓ Prevents flattening or oversimplification 1
- ✓ Ensures every relationship type (association, inheritance, aggregation) is preserved 1

Thus, this implementation adopts SCT as the mandatory architecture.<sup>1</sup>

## II. Pre-Implementation Requirements and System Setup

### A. Hardware & Environment

The environment must support:

- Running RDBMS and Neo4j simultaneously
- High I/O disk throughput (for Extract and Load phases)
- Sufficient RAM (for memory-intensive Transformation phase)
- Multi-core CPU (for parallel ETL operations)

### B. Software Stack

The selected stack aligns perfectly with techniques in the research paper.<sup>1</sup>

Component	Tool	Purpose in Migration (SCT Phase)	Source from Paper
Source RDBMS	MySQL, PostgreSQL	Hosting source data and explicit schema metadata (RDBMS) <sup>1</sup>	<sup>1</sup>
Target NoSQL	Neo4j	Target environment for property graph model <sup>1</sup>	<sup>1</sup>
Connectivity	JDBC Driver, SQL	Extracting schema metadata and data instances <sup>1</sup>	<sup>1</sup>

Transformation/Scripting	Python / Java ORM	Implementing complex schema mapping rules, data transformation <sup>2</sup>	1
Conceptual Modeling	ERD / EERD / UML	Visualizing and documenting the semantically enriched schema (C) <sup>1</sup>	1
ETL Tools	Neo4j ETL, APOC, CSV	Automated metadata extraction and high-volume data loading <sup>3</sup>	1

### III. Phase 1 — Semantic Enrichment (S → C)

This step aligns closely with the research paper's focus on metadata extraction and semantic enrichment.<sup>1</sup>

#### A. Metadata Extraction (Relational Schema Representation — RSR)

The process involves extracting explicit metadata, including tables, primary keys, foreign keys, and data types, from the source RDBMS.<sup>5</sup> This metadata is stored in an internal Relational Schema Representation (RSR) to simplify key matching and classification of relational constructs.<sup>5</sup>

#### B. Semantic Enrichment (DBRE Method)

Semantic enrichment uses Database Reverse Engineering (DBRE) techniques to infer semantics hidden in the relational schemas and data.<sup>11</sup> This is critical for generating a non-flat

target schema.<sup>1</sup>

1. **Relationship Cardinality Inference and Naming:** Analyze Foreign Keys (FKs) to determine precise cardinality (1:1, 1:N) and infer meaningful, business-oriented relationship names (e.g., MANAGES) rather than generic FK names.<sup>1</sup>
2. **Inheritance Detection:** Identify patterns such as Class Table Inheritance where related tables share an identical Primary Key structure, and the subclass PK acts as an FK to the superclass.<sup>5</sup> The system must implement pattern recognition algorithms to identify these structural commonalities.<sup>1</sup>
3. **Aggregation Identification:** Identify complex FK structures, often through analysis of cascading deletes or mandatory participation, that suggest a strong dependent lifecycle or "part-of" semantic, which must be captured explicitly.<sup>1</sup>

## C. Construction of the Conceptual Model (EER/UML)

The resulting rich, explicit Conceptual Model (CDM), documented as an EERD or UML Class Diagram, serves as the high-fidelity input for schema translation rules, ensuring a semantically strong target structure.<sup>1</sup>

## IV. Phase 2 — Schema Translation (C → Graph)

The translation dictates the precise rules for converting the enriched Conceptual Model (C) into the Property Graph Model (T).<sup>1</sup>

### A. General Mapping Rules: Entity to Node

7

1. **Entity Mapping:** Each conceptual entity maps to a unique **Node Label** (:LabelName).
2. **Row and Attribute Mapping:** Each row instance becomes a **Node**, and attributes become **properties** on that node (key: value).
3. **Key Mapping:** Technical primary keys are removed; essential **business primary keys** are mapped to **Unique Constraints** to enforce data integrity and optimize lookup performance (CREATE CONSTRAINT).<sup>7</sup> Indexes are applied to frequently searched

attributes.<sup>7</sup>

## B. Core Rule Set 1: Associative Relationship Translation

7

1. **1:N Relationships:** The foreign key column is dropped and replaced by a single, directed Edge (--)>.
2. **M:N Relationships (Join Tables):** The join table is eliminated conceptually and replaced by an Edge. If the join table contained attributes, these attributes are translated into properties of the new relationship edge.<sup>7</sup>
3. **Ternary Relationship Conversion:** The intersection table is converted into a dedicated Relationship Node to maintain connectivity and attributes, linked by three or more distinct relationships to the respective entity nodes.

## C. Core Rule Set 2: Advanced Semantic Translation

This rule set preserves the complex semantics identified in Phase 1, which are often disregarded.<sup>1</sup>

1. **Inheritance Modeling Rules:** The inheritance structure (superclass/subclass) is modeled using the **Multi-Labeling Strategy**.<sup>8</sup> The subclass node receives both the Superclass Label and its specific Subclass Label (e.g., (:Person:Employee)) by merging data from both tables onto a single node instance. This supports polymorphic queries.<sup>8</sup>
2. **Aggregation Modeling Rules:** Aggregation (strong compositional semantics) is modeled using **highly specific, specialized relationship types** (e.g., -->) to preserve the structural dependency.<sup>1</sup>

RDBMS Construct	Conceptual Inference (EERD)	Target Graph Model (Cypher/Neo4j)	Key Action
Entity Table	Entity Type	Node Label (:LabelName)	Create Node <sup>7</sup>

Row	Entity Instance	Node	Create Node <sup>7</sup>
Column (Non-Key)	Attribute	Node Property (key: value)	Set Property <sup>7</sup>
Primary Key (Business)	Unique Identifier	Unique Constraint	CREATE CONSTRAINT <sup>8</sup>
Foreign Key (1:N)	Association Relationship	Directed Relationship (-->)	Drop FK, Create Edge <sup>7</sup>
Join Table (M:N)	Association Relationship	Directed Relationship with Properties	Eliminate Table, Create Edge <sup>7</sup>
Tables with Shared PK	Inheritance (Generalization)	Node with Multiple Labels (:Superclass:Subcla ss)	Combine Entity Semantics <sup>8</sup>
Complex FK Structure	Aggregation (Composition/Part- of)	Specialized, strong relationship (-->)	Preserve Strong Semantic <sup>1</sup>
Intersection Table (3+ FKS)	Ternary Relationship	Dedicated Intersection Node	Introduce New Node Type

## V. Phase 3 — ETL: Extract, Transform, Load

### A. Extract Phase: Data Retrieval Strategy

Data extraction utilizes the **JDBC driver** to query the RDBMS.<sup>1</sup> Efficient SELECT statements retrieve data instances, which are then staged as **CSV files**, the most effective input format

for high-speed bulk import.<sup>1</sup>

## B. Transformation Phase: ORM and Data Structuring

Custom **Python or Java** scripts implement the transformation logic, executing the rules defined in Phase 2.<sup>2</sup>

1. **Data-to-Object Conversion:** Raw tabular data is mapped to intermediary objects using **ORM techniques**.<sup>1</sup>
2. **Implementing Complex Merges:** Scripts join superclass/subclass data (based on shared PKs) to create single, multi-labeled records for inheritance.<sup>8</sup> They analyze FK values to generate explicit relationship creation commands using the semantically correct names (e.g., -->).<sup>1</sup>
3. **Output Generation:** Structured CSV files are generated, separated by target component type (e.g., separate files for nodes and relationships), optimized for Neo4j's bulk loading utilities.<sup>9</sup>

## C. Load Phase: Bulk Data Import into Neo4j

High-performance loading methods are mandatory for large-scale migration.<sup>9</sup>

1. **High-Volume Bulk Load:** The **neo4j-admin database import full** utility is required for initial, large-scale imports from structured CSV files.<sup>9</sup>
2. **Medium-Volume Load:** For smaller volumes (up to 10 million records) or incremental updates, the Cypher command **LOAD CSV** is used.<sup>9</sup>
3. **API Load:** Custom scripts utilizing language libraries (e.g., Python/Java drivers, **py2neo**) can programmatically execute batched Cypher statements for highly complex or incremental updates.<sup>2</sup> The **APOC library** extends Cypher for advanced procedures during loading.<sup>9</sup>

## VI. Demonstrative Implementation Details

To ensure complete understanding and reproducibility, the following sections detail the

architectural flow, detection algorithms, and a running example of the transformation.

## A. Conceptual Architecture Flow

The complete system is based on the SCT (Source  $\rightarrow$  Conceptual  $\rightarrow$  Target) methodology, requiring a sequence of specialized components:

1. **RDBMS (Source):** Stores source data and explicit schema metadata (tables, FKs, PKs).<sup>1</sup>
2. **Schema Analyzer (JDBC):** Connects to the RDBMS using the JDBC Driver<sup>1</sup> to extract metadata and build the **Relational Schema Representation (RSR)**.<sup>5</sup>
3. **Conceptual Model Builder (DBRE):** Performs **Semantic Enrichment** on the RSR, using techniques like Database Reverse Engineering (DBRE)<sup>11</sup> to infer implicit semantics (e.g., **Inheritance** and **Aggregation** patterns).<sup>1</sup> This results in the **Canonical Data Model (CDM)** / EERD.<sup>5</sup>
4. **Transformation Engine (ORM/Scripting):** Uses the CDM rules to drive the process. It extracts data instances from the RDBMS, applies Object-Relational Mapping (ORM)<sup>1</sup>, performs complex data merges (e.g., inheritance joining), and generates staged **CSV files** for nodes and relationships.<sup>9</sup>
5. **Neo4j Loader (ETL Tools):** Ingests the CSV files into the **Neo4j Graph Database (Target)** using high-volume utilities like neo4j-admin import full.<sup>9</sup>

## B. Algorithmic Pseudocode for Semantic Detection and Mapping

### Algorithm 1: Inference and Merging for Inheritance (Generalization)

```
// INPUT: RSR containing Table MetaData (T1, T2)
// OUTPUT: Transformation Rule to merge T1 and T2 data into one Multi-Labeled Node

FUNCTION InferAndMergeInheritance(Table T1, Table T2):
    // Check for shared PK structure and FK relationship
    IF T1.PrimaryKey == T2.PrimaryKey AND T2.PrimaryKey IS ForeignKeyTo T1.PrimaryKey:
        Superclass = T1
        Subclass = T2
```

```

// Apply Multi-Labeling Strategy
TransformationRule = {
    SourceTables:,
    TargetLabel: Superclass.Name + ":" + Subclass.Name,
    MergeAction: JOIN_ON_PRIMARY_KEY,
    Output: Single_Record_With_MultiLabels
}
RETURN TransformationRule
END IF
RETURN NULL

// Implementation requires script to JOIN Superclass and Subclass data rows based on their
// shared key
// and assign the combined labels (e.g., (:Person:Employee)) to the resulting node instance.

```

Algorithm 2: Mapping M:N Join Tables to Relationships <sup>7</sup>

```

// INPUT: RSR containing Join Table J and two related Entity Tables A, B
// OUTPUT: Cypher Generation Rule for Relationship creation

FUNCTION MapJoinTableToRelationship(JoinTable J, Entity A, Entity B):

    // Identify the semantic relationship name (derived in Phase 1)
    RelationshipName = J.InferredName // e.g., 'ENROLLED_IN'

    // Determine if J holds properties (columns other than the two FKs)
    RelationshipProperties = J.Attributes.EXCLUDE(A.ForeignKey, B.ForeignKey)

    // Create the relationship structure
    IF RelationshipProperties IS NOT EMPTY:
        // Attributes become properties of the relationship
        Rule = {
            Source: A.Name,
            Target: B.Name,
            RelationshipType: RelationshipName,
            Properties: RelationshipProperties
        }
    ELSE:

```

```

// Simple relationship (no properties)
Rule = {
    Source: A.Name,
    Target: B.Name,
    RelationshipType: RelationshipName,
    Properties:
}
END IF

// Action in Transform Phase: Eliminate the table and generate relationship creation
commands
RETURN Rule

```

### C. Running Example: Student-Course-Grade Model

This example demonstrates the translation path for key constructs from the Source RDBMS to the Target Graph Schema.

Source RDBMS Construct	Conceptual Enrichment (Phase 1)	Target Graph Schema (Phase 2)	Transformation Summary
<b>Table:</b> Course (ID:PK, Title, DeptID:FK)	Entity Type (Course) + 1:N Association to Department	<b>Node Label:</b> (:Course) <sup>7</sup>	ID (FK source) is dropped; Node is created.
<b>Table:</b> Department (DeptID:PK, Name)	Entity Type (Department)	<b>Node Label:</b> (:Department) <sup>7</sup>	Node is created with Name property.
<b>FK:</b> Course.DeptID → Department.DeptID <sup>1</sup>	Inferred Relationship: <b>OFFERED_BY</b> (1:N)	<b>Relationship Edge:</b> (:Course)-->(:Department) <sup>7</sup>	Foreign key column is removed and replaced by a directed, semantically named relationship. <sup>7</sup>

<b>Table:</b> Takes (StudentID:FK, CourseID:FK, <b>Grade</b> )	M:N Join Table with Property (Grade) <sup>7</sup>	<b>Relationship Edge with Property:</b> (:Student)-->(:Course) <sup>7</sup>	Join table is eliminated; its non-key attribute (Grade) becomes a property of the new relationship edge. <sup>7</sup>
---	--	--	--

## VII. Verification, Validation, and Performance Optimization

The project concludes with rigorous validation to confirm that the migration was successful in both preserving data integrity and delivering architectural advantages, specifically query performance gains.<sup>10</sup>

### A. Data Instance and Semantic Integrity Validation

Verification must confirm that all data instances were transferred and that the complex semantic models were correctly realized.<sup>10</sup>

1. **Data Integrity (Count Verification):** Compare the number of rows extracted from the source RDBMS tables with the total number of nodes created for the corresponding labels in Neo4j.<sup>10</sup> The number of edges created must precisely correspond to the number of foreign key instances processed.<sup>10</sup>
2. **Constraint and Semantic Fidelity Check:** Verify that all Unique Constraints defined on the Neo4j nodes (corresponding to business keys) are correctly enforced.<sup>7</sup> Specific Cypher queries must be run to validate complex semantic translations, such as confirming the multi-labeling inheritance model is functional.<sup>1</sup>

### B. Performance Validation and Query Translation

A key objective is the reduction of query cost, particularly minimizing join time.<sup>1</sup>

1. **SQL Query \$\rightarrow\$ Cypher Query Translation:** Identify a benchmark set of complex, multi-join SQL queries from the RDBMS. Translate these costly queries into equivalent, path-finding Cypher traversal patterns.<sup>13</sup>

2. **Performance Benchmarking:** Execute both the original SQL queries and their Cypher counterparts and compare their average execution times (latency).<sup>10</sup> A successful project demonstrates that the Cypher traversal queries execute significantly faster than the complex, depth-dependent SQL joins.<sup>1</sup>

## C. Iterative Refinement and Graph Model Optimization

Refine the graph model based on performance benchmarking.<sup>8</sup> Ensure all frequently used lookup properties are correctly indexed for fast traversal entry points.<sup>7</sup> Optimization may include refining relationship types or adjusting denormalization strategies to ensure optimal performance.<sup>10</sup>

Validation Phase	Metric/Test Goal	Data to Compare (RDBMS vs. Neo4j)	Success Criterion
Data Integrity	Data Count Preservation	Row Counts per RDBMS Table vs. Node Counts per Neo4j Label	Counts must match exactly (or match defined merge criteria). <sup>10</sup>
Semantic Fidelity	Constraint Enforcement	RDBMS Primary/Foreign Keys vs. Neo4j Constraints/Indexe s	Unique constraints are enforced; derived relationship types are queryable and accurate. <sup>7</sup>
Functional Equivalence	Query Latency	Execution time of complex multi-join SQL queries vs. equivalent Cypher traversal queries. <sup>1</sup>	Cypher query execution time is significantly reduced, demonstrating architectural benefit. <sup>1</sup>

Model Optimization	Index Usage	Frequency of index hits and cache performance in Neo4j. <sup>7</sup>	All critical lookup properties are indexed; traversal paths are optimized. <sup>10</sup>
--------------------	-------------	--	--

## Works cited

1. Data Migration from Conventional Databases into NoSQL\_ Methods and Techniques.pdf
2. Migrating mysql data to neo4j database - Stack Overflow, accessed November 19, 2025,  
<https://stackoverflow.com/questions/24882715/migrating-mysql-data-to-neo4j-database>
3. Neo4j ETL Tool - Interactive Relational Database Data Import, accessed November 19, 2025, <https://neo4j.com/labs/etl-tool/>
4. Import: RDBMS to graph - Getting Started - Neo4j, accessed November 19, 2025, <https://neo4j.com/docs/getting-started/data-import/relational-to-graph-import/>
5. Semantic Enrichment: The First Phase of Relational Database Migration - ResearchGate, accessed November 19, 2025,  
[https://www.researchgate.net/publication/274719624\\_Semantic\\_Enrichment\\_The\\_First\\_Phase\\_of\\_Relational\\_Database\\_Migration](https://www.researchgate.net/publication/274719624_Semantic_Enrichment_The_First_Phase_of_Relational_Database_Migration)
6. Database Schema Design Using Entity-Relationship Approach, accessed November 19, 2025, <https://www.comp.nus.edu.sg/~lingtw/cs4221/er.pdf>
7. Modeling: relational to graph - Getting Started - Neo4j, accessed November 19, 2025,  
<https://neo4j.com/docs/getting-started/data-modeling/relational-to-graph-modeling/>
8. A Rule-based Conversion of an EER Schema to Neo4j Schema ..., accessed November 19, 2025,  
[https://www.researchgate.net/publication/357011565\\_A\\_Rule-based\\_Conversion\\_of\\_an\\_EER\\_Schema\\_to\\_Neo4j\\_Schema\\_Constraints](https://www.researchgate.net/publication/357011565_A_Rule-based_Conversion_of_an_EER_Schema_to_Neo4j_Schema_Constraints)
9. Import your data into Neo4j - Getting Started, accessed November 19, 2025, <https://neo4j.com/docs/getting-started/data-import/>
10. Migrate from Relational Database to Graph Database - FalkorDB, accessed November 19, 2025,  
<https://www.falkordb.com/blog/relational-database-to-graph-database-2/>
11. OWL Ontology Extraction from Relational Databases via Database Reverse Engineering - Journal of Software, accessed November 19, 2025, <https://www.jsoftware.us/vol8/jsw0811-11.pdf>
12. RDBMS & Graphs: Relational vs. Graph Data Modeling - Neo4j, accessed November 19, 2025, <https://neo4j.com/blog/developer/rdbms-vs-graph-data-modeling/>
13. Cypher Query VS SQL. Databases have become an integral part... | by Nimra Tahir | Medium, accessed November 19, 2025,

<https://medium.com/@nimratahir1212/cypher-query-vs-sql-718ccb8e8d9d>